

# Codsoft Task 2 - Credit Card Fraud Detection

For this task, we'll be using Logistic Regression, Decision Trees or Random Forests to classify transactions as fraudulent or legitimate

## Importing required libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import plotly.express as px
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: train = pd.read_csv('./Dataset/fraudTrain.csv', index_col=0)
test = pd.read_csv('./Dataset/fraudTest.csv', index_col=0)
```

## Lets Analyze Training data

```
In [3]: train.head()
```

Out[3]:	trans_date_trans_time	cc_num	merchant	category	amt	first
<b>0</b>	2019-01-01 00:00:18	2703186189652095	fraud_Rippin, Kub and Mann	misc_net	4.97	Jennifer
<b>1</b>	2019-01-01 00:00:44	630423337322	fraud_Heller, Gutmann and Zieme	grocery_pos	107.23	Stephanie
<b>2</b>	2019-01-01 00:00:51	38859492057661	fraud_Lind- Buckridge	entertainment	220.11	Edward
<b>3</b>	2019-01-01 00:01:16	3534093764340240	fraud_Kutch, Hermiston and Farrell	gas_transport	45.00	Jeremy
<b>4</b>	2019-01-01 00:03:06	375534208663984	fraud_Keeling- Crist	misc_pos	41.96	Tyler

5 rows × 22 columns

In [4]: `train.info()`

```

<class 'pandas.core.frame.DataFrame'>
Index: 1296675 entries, 0 to 1296674
Data columns (total 22 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   trans_date_trans_time                 1296675 non-null object
1   cc_num                               1296675 non-null int64
2   merchant                             1296675 non-null object
3   category                             1296675 non-null object
4   amt                                   1296675 non-null float64
5   first                                1296675 non-null object
6   last                                  1296675 non-null object
7   gender                               1296675 non-null object
8   street                               1296675 non-null object
9   city                                  1296675 non-null object
10  state                                1296675 non-null object
11  zip                                   1296675 non-null int64
12  lat                                   1296675 non-null float64
13  long                                  1296675 non-null float64
14  city_pop                              1296675 non-null int64
15  job                                   1296675 non-null object
16  dob                                   1296675 non-null object
17  trans_num                             1296675 non-null object
18  unix_time                             1296675 non-null int64
19  merch_lat                             1296675 non-null float64
20  merch_long                            1296675 non-null float64
21  is_fraud                              1296675 non-null int64
dtypes: float64(5), int64(5), object(12)
memory usage: 227.5+ MB

```

```
In [5]: train.shape
```

```
Out[5]: (1296675, 22)
```

```
In [6]: train.describe()
```

```
Out[6]:
```

	cc_num	amt	zip	lat	long	city_p
<b>count</b>	1.296675e+06	1.296675e+06	1.296675e+06	1.296675e+06	1.296675e+06	1.296675e+
<b>mean</b>	4.171920e+17	7.035104e+01	4.880067e+04	3.853762e+01	-9.022634e+01	8.882444e+
<b>std</b>	1.308806e+18	1.603160e+02	2.689322e+04	5.075808e+00	1.375908e+01	3.019564e+
<b>min</b>	6.041621e+10	1.000000e+00	1.257000e+03	2.002710e+01	-1.656723e+02	2.300000e+
<b>25%</b>	1.800429e+14	9.650000e+00	2.623700e+04	3.462050e+01	-9.679800e+01	7.430000e+
<b>50%</b>	3.521417e+15	4.752000e+01	4.817400e+04	3.935430e+01	-8.747690e+01	2.456000e+
<b>75%</b>	4.642255e+15	8.314000e+01	7.204200e+04	4.194040e+01	-8.015800e+01	2.032800e+
<b>max</b>	4.992346e+18	2.894890e+04	9.978300e+04	6.669330e+01	-6.795030e+01	2.906700e+

```
In [7]: train.isnull().sum()
```

```
Out[7]: trans_date_trans_time    0
        cc_num                  0
        merchant                 0
        category                 0
        amt                     0
        first                    0
        last                     0
        gender                   0
        street                   0
        city                     0
        state                    0
        zip                      0
        lat                      0
        long                     0
        city_pop                 0
        job                      0
        dob                      0
        trans_num                0
        unix_time                0
        merch_lat                0
        merch_long               0
        is_fraud                 0
        dtype: int64
```

```
In [8]: train.duplicated().sum()
```

```
Out[8]: 0
```

**From the above we can clearly see that there are no duplicates and null values in the training dataset.**

```
In [9]: train['is_fraud'].value_counts()
```

```
Out[9]: is_fraud
0      1289169
1         7506
Name: count, dtype: int64
```

**We can see that the training data is imbalanced**

**Lets Analyze the test data**

```
In [10]: test.head()
```

Out[10]:

	trans_date_trans_time	cc_num	merchant	category	amt	first
0	2020-06-21 12:14:25	2291163933867244	fraud_Kirlin and Sons	personal_care	2.86	Jeff
1	2020-06-21 12:14:33	3573030041201292	fraud_Sporer-Keebler	personal_care	29.84	Joanne
2	2020-06-21 12:14:53	3598215285024754	fraud_Swaniawski, Nitzsche and Welch	health_fitness	41.28	Ashley
3	2020-06-21 12:15:15	3591919803438423	fraud_Haley Group	misc_pos	60.05	Brian
4	2020-06-21 12:15:17	3526826139003047	fraud_Johnston-Casper	travel	3.19	Nathan

5 rows × 22 columns

In [11]: `test.info()`

```

<class 'pandas.core.frame.DataFrame'>
Index: 555719 entries, 0 to 555718
Data columns (total 22 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   trans_date_trans_time                 555719 non-null object
1   cc_num                               555719 non-null int64
2   merchant                             555719 non-null object
3   category                             555719 non-null object
4   amt                                   555719 non-null float64
5   first                                555719 non-null object
6   last                                  555719 non-null object
7   gender                               555719 non-null object
8   street                               555719 non-null object
9   city                                  555719 non-null object
10  state                                555719 non-null object
11  zip                                   555719 non-null int64
12  lat                                   555719 non-null float64
13  long                                  555719 non-null float64
14  city_pop                             555719 non-null int64
15  job                                   555719 non-null object
16  dob                                   555719 non-null object
17  trans_num                            555719 non-null object
18  unix_time                            555719 non-null int64
19  merch_lat                            555719 non-null float64
20  merch_long                           555719 non-null float64
21  is_fraud                             555719 non-null int64
dtypes: float64(5), int64(5), object(12)
memory usage: 97.5+ MB

```

```
In [12]: test.shape
```

```
Out[12]: (555719, 22)
```

```
In [13]: test.describe()
```

```
Out[13]:
```

	cc_num	amt	zip	lat	long	city
<b>count</b>	5.557190e+05	555719.000000	555719.000000	555719.000000	555719.000000	5.557190e+05
<b>mean</b>	4.178387e+17	69.392810	48842.628015	38.543253	-90.231325	8.822189e+06
<b>std</b>	1.309837e+18	156.745941	26855.283328	5.061336	13.721780	3.003909e+07
<b>min</b>	6.041621e+10	1.000000	1257.000000	20.027100	-165.672300	2.300000e+06
<b>25%</b>	1.800429e+14	9.630000	26292.000000	34.668900	-96.798000	7.410000e+06
<b>50%</b>	3.521417e+15	47.290000	48174.000000	39.371600	-87.476900	2.408000e+07
<b>75%</b>	4.635331e+15	83.010000	72011.000000	41.894800	-80.175200	1.968500e+07
<b>max</b>	4.992346e+18	22768.110000	99921.000000	65.689900	-67.950300	2.906700e+07

```
In [14]: test.isnull().sum()
```


```
Out[14]: trans_date_trans_time    0
         cc_num                  0
         merchant                 0
         category                 0
         amt                     0
         first                   0
         last                    0
         gender                  0
         street                  0
         city                    0
         state                   0
         zip                     0
         lat                     0
         long                    0
         city_pop                0
         job                     0
         dob                     0
         trans_num               0
         unix_time               0
         merch_lat               0
         merch_long              0
         is_fraud                0
         dtype: int64
```

```
In [15]: test.duplicated().sum()
```

```
Out[15]: 0
```


## Visualization

```
In [16]: cat_his = px.histogram(data_frame=train,x='category',color='is_fraud');
         cat_his.show();
```



```
In [17]: gen_his = px.histogram(data_frame=train,x='gender',color='is_fraud');  
gen_his.show()
```





```
In [18]: train.columns
```

```
Out[18]: Index(['trans_date_trans_time', 'cc_num', 'merchant', 'category', 'amt',  
               'first', 'last', 'gender', 'street', 'city', 'state', 'zip', 'lat',  
               'long', 'city_pop', 'job', 'dob', 'trans_num', 'unix_time', 'merch_lat',  
               'merch_long', 'is_fraud'],  
              dtype='object')
```

```
In [19]: test.columns
```

```
Out[19]: Index(['trans_date_trans_time', 'cc_num', 'merchant', 'category', 'amt',  
               'first', 'last', 'gender', 'street', 'city', 'state', 'zip', 'lat',  
               'long', 'city_pop', 'job', 'dob', 'trans_num', 'unix_time', 'merch_lat',  
               'merch_long', 'is_fraud'],  
              dtype='object')
```

## Data Processing & Cleaning

```
In [20]: #Date of birth --> Age of customer
```

```
train['dob'] = pd.to_datetime(train['dob'],format='mixed')  
test['dob'] = pd.to_datetime(test['dob'],format='mixed')
```

```

train['trans_date_trans_time'] = pd.to_datetime(train['trans_date_trans_time'],format
test['trans_date_trans_time'] = pd.to_datetime(test['trans_date_trans_time'],format

train['age'] = (train['trans_date_trans_time'].dt.year - train['dob'].dt.year).astype
test['age'] = (test['trans_date_trans_time'].dt.year - test['dob'].dt.year).astype(

train.drop(columns='dob',inplace=True)
test.drop(columns='dob',inplace=True)

```

```

In [21]: #cleaning merchant column
train['merchant'] = train['merchant'].apply(lambda x : x.replace('fraud_', ''))
test['merchant'] = test['merchant'].apply(lambda x : x.replace('fraud_', ''))

```

```

In [22]: #Converting gender to binary classification
train = pd.get_dummies(train,columns=['gender'],drop_first=True)
test = pd.get_dummies(test,columns=['gender'],drop_first=True)

```

```

In [23]: train_numerical_columns = train.select_dtypes(include=['int64', 'float64']).columns
train_numerical_columns

```

```

Out[23]: Index(['cc_num', 'amt', 'zip', 'lat', 'long', 'city_pop', 'unix_time',
               'merch_lat', 'merch_long', 'is_fraud'],
              dtype='object')

```

```

In [24]: test_numerical_columns = test.select_dtypes(include=['int64', 'float64']).columns
test_numerical_columns

```

```

Out[24]: Index(['cc_num', 'amt', 'zip', 'lat', 'long', 'city_pop', 'unix_time',
               'merch_lat', 'merch_long', 'is_fraud'],
              dtype='object')

```

```

In [25]: train.drop(
          columns = ['cc_num','trans_date_trans_time','first', 'last','street','state',
                    'zip', 'lat','long', 'city_pop','unix_time', 'merch_lat','merch_long'
          test.drop(
          columns = ['cc_num','trans_date_trans_time','first', 'last','street','state',
                    'zip', 'lat','long', 'city_pop','unix_time', 'merch_lat','merch_long'

```

```

In [26]: train = train[['merchant','category','city','job','gender_M','age','amt','is_fraud']
train.head()

```

Out[26]:

	merchant	category	city	job	gender_M	age	amt	is_fraud
0	Rippin, Kub and Mann	misc_net	Moravian Falls	Psychologist, counselling	False	31	4.97	0
1	Heller, Gutmann and Zieme	grocery_pos	Orient	Special educational needs teacher	False	41	107.23	0
2	Lind-Buckridge	entertainment	Malad City	Nature conservation officer	True	57	220.11	0
3	Kutch, Hermiston and Farrell	gas_transport	Boulder	Patent attorney	True	52	45.00	0
4	Keeling-Crist	misc_pos	Doe Hill	Dance movement psychotherapist	True	33	41.96	0

In [27]: `test = test[['merchant', 'category', 'city', 'job', 'gender_M', 'age', 'amt', 'is_fraud']]`  
`test.head()`

Out[27]:

	merchant	category	city	job	gender_M	age	amt	is_fraud
0	Kirlin and Sons	personal_care	Columbia	Mechanical engineer	True	52	2.86	0
1	Sporer-Keebler	personal_care	Altonah	Sales professional, IT	False	30	29.84	0
2	Swaniawski, Nitzsche and Welch	health_fitness	Bellmore	Librarian, public	False	50	41.28	0
3	Haley Group	misc_pos	Titusville	Set designer	True	33	60.05	0
4	Johnston-Casper	travel	Falmouth	Furniture designer	True	65	3.19	0

## Now we wil encode the String columns

In [28]: `# We will use Weight of Evidence encoder for this task`  
`from category_encoders import WOEEncoder`

In [29]: `encoder = WOEEncoder(cols=['city', 'job', 'merchant', 'category'])`  
`train_transformed = encoder.fit_transform(train[['city', 'job', 'merchant', 'category']])`  
`test_transformed = encoder.transform(test[['city', 'job', 'merchant', 'category']])`

```
# Replace the original columns with the transformed values
train[['city', 'job', 'merchant', 'category']] = train_transformed
test[['city', 'job', 'merchant', 'category']] = test_transformed
```

In [30]: `train.head(3)`

```
Out[30]:
```

	merchant	category	city	job	gender_M	age	amt	is_fraud
0	0.959326	0.924914	-2.469513	-1.080186	False	31	4.97	0
1	0.663187	0.898799	-3.027790	-0.904144	False	41	107.23	0
2	-0.790166	-0.847622	-1.076791	1.120434	True	57	220.11	0

In [31]: `test.head(3)`

```
Out[31]:
```

	merchant	category	city	job	gender_M	age	amt	is_fraud
0	-1.259443	-0.869588	0.364725	0.430148	True	52	2.86	0
1	-0.569596	-0.869588	0.215637	0.200487	False	30	29.84	0
2	-1.202771	-1.315531	-0.626823	-0.637068	False	50	41.28	0

We can see that the data is not standardized. Lets Standardize

In [32]: `from sklearn.preprocessing import StandardScaler`

In [33]: `scaler = StandardScaler()`

In [34]: `X_train = train.drop('is_fraud',axis=1)`  
`y_train = train['is_fraud']`  
`X_test = test.drop('is_fraud',axis=1)`  
`y_test = test['is_fraud']`

In [35]: `X_train = scaler.fit_transform(X_train)`  
`X_test = scaler.transform(X_test)`

Now we have Standardized the Train & Test Data

## Model Building

First we will proceed with Models without sampling the data. Then we will sample the data and check the accuracies before and after Sampling

In [36]: `from sklearn.linear_model import LogisticRegression`  
`from sklearn.tree import DecisionTreeClassifier`  
`from sklearn.ensemble import RandomForestClassifier`

```
from sklearn.metrics import *
```

```
In [37]: models = [LogisticRegression(),DecisionTreeClassifier()]  
# Removed RandomForestClassifier() model as it is taking very long time for predict
```

```
In [38]: def model_prediction(X_train,y_train,X_test,y_test):  
    for model in models:  
        model.fit(X_train,y_train)  
        y_prediction = model.predict(X_test)  
  
        print(f"{model} Model")  
        print("\nAccuracy Score :",accuracy_score(y_test,y_prediction))  
        print("Precision :",precision_score(y_test,y_prediction))  
        print("Recall Score :",recall_score(y_test,y_prediction))  
        print("F1 Score :",f1_score(y_test,y_prediction))  
        print(".....\n")
```

```
In [39]: model_prediction(X_train,y_train,X_test,y_test)
```

LogisticRegression() Model

Accuracy Score : 0.9956704737466238

Precision : 0.03225806451612903

Recall Score : 0.004195804195804196

F1 Score : 0.007425742574257425

.....

DecisionTreeClassifier() Model

Accuracy Score : 0.9951612235680263

Precision : 0.3353510895883777

Recall Score : 0.25827505827505826

F1 Score : 0.29180932314985514

.....

## Lets Apply Sampling and check the Accuracies

```
In [40]: from imblearn.over_sampling import SMOTE
```

```
In [41]: smote = SMOTE()  
X_train,y_train = smote.fit_resample(X_train,y_train)
```

```
In [42]: model_prediction(X_train,y_train,X_test,y_test)
```

LogisticRegression() Model

Accuracy Score : 0.9121498455154493

Precision : 0.02229136388758111

Recall Score : 0.5076923076923077

F1 Score : 0.04270755715910428

.....

DecisionTreeClassifier() Model

Accuracy Score : 0.9938026232682344

Precision : 0.19776640297812936

Recall Score : 0.19813519813519814

F1 Score : 0.19795062878435024

.....

**From the above, We can clearly see that the accuracies, precisions, recall & F1 score are very low. Now we try with dropping String columns from the train and test dataset.**

```
In [43]: fraudTrain_df = pd.read_csv('./Dataset/fraudTrain.csv',index_col=0)
        fraudTest_df = pd.read_csv('./Dataset/fraudTrain.csv',index_col=0)
```

```
In [44]: # Remove string columns
        numeric_fraudTrain_df = fraudTrain_df.select_dtypes(exclude=['object'])
        numeric_fraudTest_df = fraudTest_df.select_dtypes(exclude=['object'])
```

```
In [45]: X_train = train.drop('is_fraud',axis=1)
        y_train = train['is_fraud']

        X_test = test.drop('is_fraud',axis=1)
        y_test = test['is_fraud']
```

```
In [46]: # Standardize the features
        scaler = StandardScaler()
        X_train_scaled = scaler.fit_transform(X_train)
        X_test_scaled = scaler.transform(X_test)
```

```
In [47]: # Initialize classifiers
        log_reg = LogisticRegression()
        dec_tree = DecisionTreeClassifier()
```

```
In [48]: # Train classifiers
        log_reg.fit(X_train_scaled, y_train)
        dec_tree.fit(X_train, y_train)
```

```
Out[48]: ▾ DecisionTreeClassifier
        DecisionTreeClassifier()
```

```
In [49]: # Make predictions
        y_pred_log_reg = log_reg.predict(X_test_scaled)
```

```
y_pred_dec_tree = dec_tree.predict(X_test)
```

```
In [50]: # Calculate metrics
metrics = {}
metrics['Logistic Regression'] = {
    'Precision': precision_score(y_test, y_pred_log_reg),
    'Recall': recall_score(y_test, y_pred_log_reg),
    'Accuracy': accuracy_score(y_test, y_pred_log_reg),
    'F1 Score': f1_score(y_test, y_pred_log_reg)
}
metrics['Decision Tree'] = {
    'Precision': precision_score(y_test, y_pred_dec_tree),
    'Recall': recall_score(y_test, y_pred_dec_tree),
    'Accuracy': accuracy_score(y_test, y_pred_dec_tree),
    'F1 Score': f1_score(y_test, y_pred_dec_tree)
}
```

```
In [51]: # Print the metrics for each model
for model, scores in metrics.items():
    print(f'{model} metrics:')
    print('\n')
    for score_name, score_value in scores.items():
        print(f'{score_name}: {score_value:.4f}')
    print('.....')
```

Logistic Regression metrics:

Precision: 0.0323  
Recall: 0.0042  
Accuracy: 0.9957  
F1 Score: 0.0074  
.....  
Decision Tree metrics:

Precision: 0.3370  
Recall: 0.2555  
Accuracy: 0.9952  
F1 Score: 0.2906  
.....